




YouTube Comment Toxicity Detector

Classification using Weighted SVM and DistilBERT.

Supervised by: **Dr. Mohamed Badwei**



Our Team



Mohamed Younis



Abdelrahman Omar



Jovanie Hanie

Problem Statement

1. The Content Moderation Crisis:


- Manual moderation of billions of daily YouTube comments is physically and economically impossible.
- Automated systems are essential to maintain safe digital environments and prevent the spread of hate speech.

2. Extreme Class Imbalance:

- The vast majority of comments in the dataset (over 143,000 out of 159,571) are non-toxic.
- Critical offensive categories, such as "Threat" and "Identity Hate," represent only a tiny fraction of the data, making it difficult for standard models to recognize them.

3. Multi-Label Complexity:

- Comments often belong to multiple categories simultaneously (e.g., "Obscene" and "Insult"), requiring a model that can predict six overlapping labels with high precision.




Dataset & The Multi-Label Challenge

- **Source:** Jigsaw Toxic Comment Classification Challenge dataset (provided by Google/Jigsaw).
- **Data Volume:** 159,571 training samples.
- **Target Labels:** The project addresses six non-exclusive categories: Toxic, Severe Toxic, Obscene, Threat, Insult, and Identity Hate.
- **Critical Problem:** Significant Class Imbalance. The vast majority of comments are "clean," while offensive categories like "Threat" and "Identity Hate" represent a tiny fraction of the data, necessitating specialized handling.



Dataset



Preprocessing & Vectorization Pipeline

Text Normalization

Utilized the SpaCy library for advanced Lemmatization and removal of non-essential components like stopwords and URLs.

Statistical Path

Implemented a TF-IDF Vectorizer limited to the top 10,000 features to convert text into numerical weights based on word frequency.


Neural Path

Employed the DistilBERT Tokenizer using subword tokenization with a maximum sequence length of 128 tokens to capture contextual meaning.



Baseline Architecture: Weighted Linear

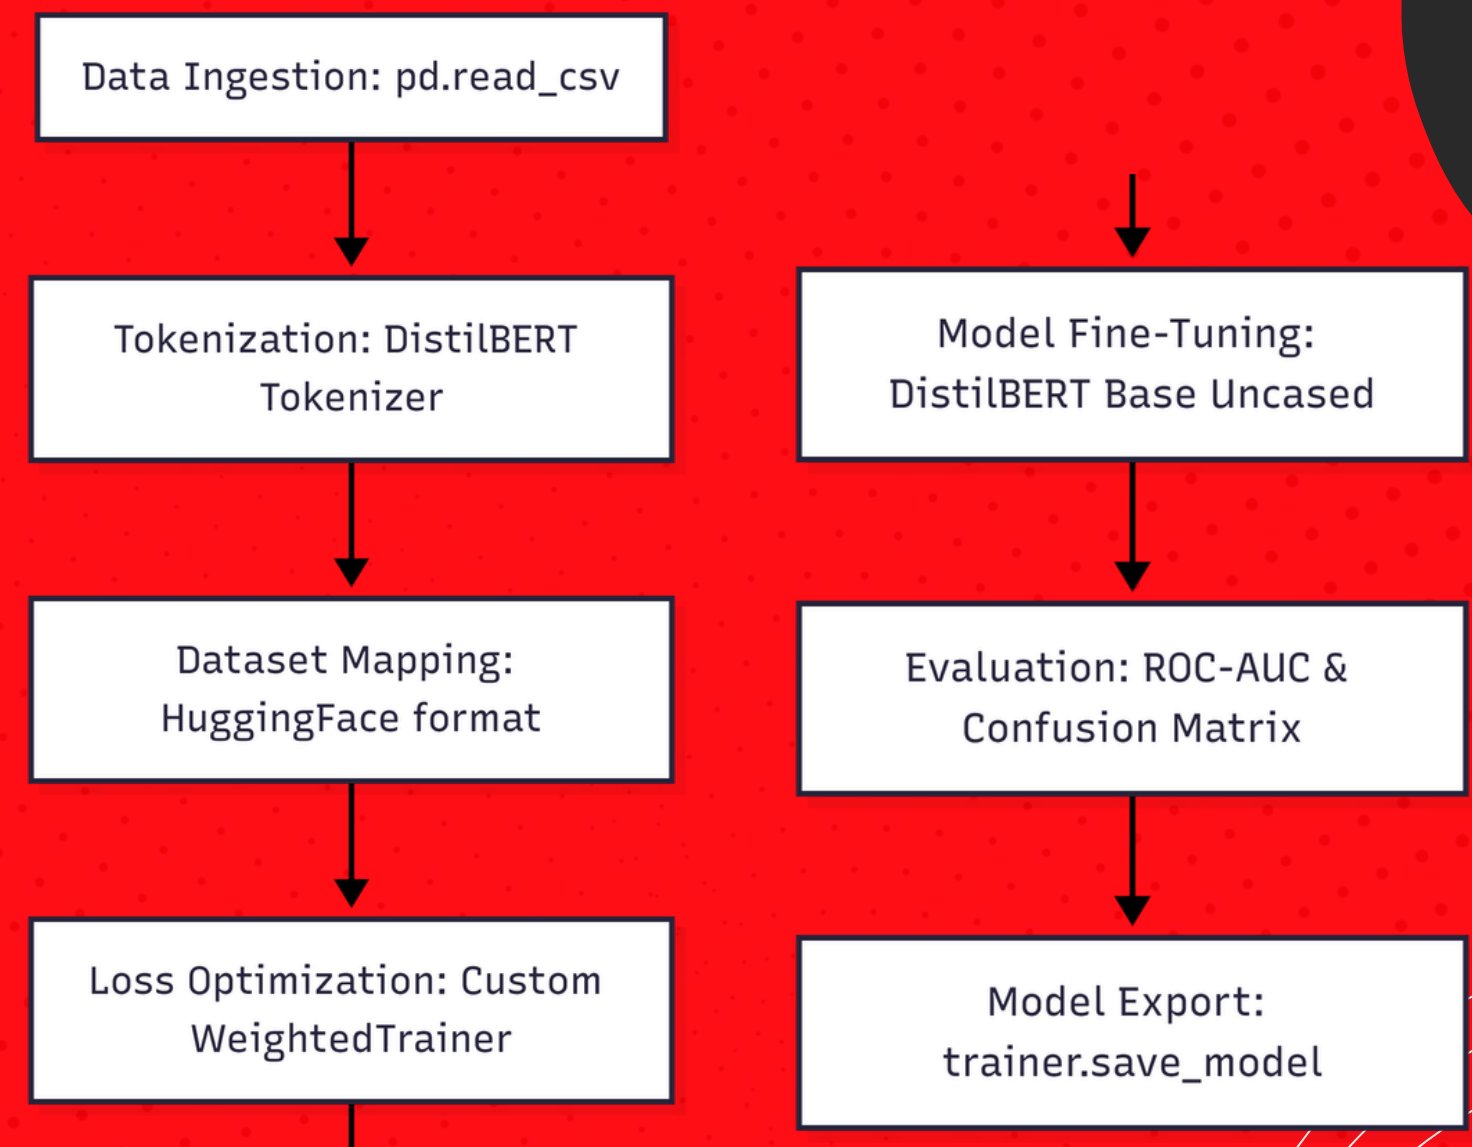
- **Approach:** Utilized a Linear regression model wrapped in a MultiOutputClassifier to handle multiple target labels simultaneously.
- **Imbalance Mitigation:** Applied the `class_weight='balanced'` parameter to automatically adjust weights inversely proportional to class frequencies.
- **Rationale:** Linear Regression was selected as the baseline because it is computationally efficient, fast for inference, and highly effective when paired with TF-IDF for large-scale text data.



Advanced Architecture: DistilBERT

Fine-tuning

- **Model Selection:** distilbert-base-uncased. It is 40% smaller and 60% faster than BERT while retaining 97% of its language understanding performance.
- **Custom Training Logic:** Developed a specialized WeightedTrainer that incorporates pos_weight within the BCEWithLogitsLoss function. This forces the model to prioritize learning from rare, high-severity labels.
- **Training Dynamics:** Used a learning rate of $2e-5$ with a Cosine Learning Rate Scheduler and a 10% warmup phase to ensure training stability and prevent overfitting.



Conclusion & Future Scope

- 📁 **Conclusion:** The project successfully developed a robust detector where DistilBERT outperformed traditional models by capturing semantic intent rather than just word frequency.
- 📁 **Deployment Potential:** The model is ready for integration with the YouTube Data API and a Streamlit dashboard for real-time comment moderation.
- 📁 **Future Directions:** Further research will explore larger models like RoBERTa and the use of Data Augmentation techniques to generate more synthetic samples for the rarest labels.

Thank You 